

Build-Your-Own SQL Database Agent (From Scratch)

Vizuara's LLM Production and Deployment

Max team size: 2 — Deadline: November 10, 2025 (23:59 IST)

1. Overview & Learning Objectives

You will build your own agent library entirely from scratch. This assignment focuses on designing and implementing a lightweight agent that interprets natural language and interacts with a SQL database using the **ReAct** (Reasoning + Acting) **Thought → Action → Observation** loop. No agent frameworks may be used.

By completing this assignment, you will:

- Implement a ReAct (Thought–Action–Observation) loop from scratch.
- Design and integrate tools for LLM-based agents.
- Create concise prompts that induce structured reasoning.
- Handle errors, timeouts, and edge cases in agent systems.
- Understand foundational ideas behind libraries such as LangChain and SmolAgents.

2. What You Will Build

A single-purpose, read-only **SQL Database Agent** that can:

- Explore database schemas (list and describe tables).
- Answer natural-language questions about the data.
- Execute *safe* SQL SELECT-only queries.
- Explain its reasoning step-by-step via auditable traces.

3. Technical Requirements

3.1 Constraints

- **Code Length:** 350–400 lines of Python (excluding comments and docstrings). This is not a strict requirement, but the focus will be on having the code as small as possible.
- **Dependencies:** One LLM client library (*OpenAI or Anthropic*), `sqlite3` (stdlib). Optional: `tabulate`.
- **Prohibited:** No agent frameworks (e.g., LangChain, SmolAgents, crew libraries).
- **Database:** SQLite (.db) read-only. Your agent must not perform mutations.

3.2 Required Components

A. Core Agent Class (120–150 lines)

- Initialize with LLM client, DB connection, and tool registry.
- Implement a bounded-step **ReAct** loop.
- Parse LLM outputs to extract **THOUGHT**, **ACTION**{json}, and parameters.
- Maintain conversation history and produce a **FINAL_ANSWER**.

B. Tool System (80–100 lines)

- Base Tool interface (name, description, parameters, call).
- Required tools:
 - `list_tables()`: list all tables.
 - `describe_table(table_name)`: schema (columns + types, and row count if convenient).
 - `query_database(query)`: SELECT-only execution with safe validation.
- Standardized tool descriptions must be surfaced to the LLM.

C. Prompt Engineering (50–70 lines)

- A concise *system prompt* enforcing ReAct formatting.
- 1–2 minimal few-shot traces (keep them short; no full solutions).
- Clear tool documentation block (names, args, returns).
- Output parsing template for **THOUGHT**, **ACTION**, **OBSERVATION**, **FINAL_ANSWER**.

D. Utilities (50–80 lines)

- SQL validation: SELECT-only; prohibit INSERT/UPDATE/DELETE/DROP/ALTER, subqueries optional to forbid.
- Result formatting (e.g., small tables to readable text).
- Error handling and retry logic; bounded time/row caps (e.g., `LIMIT`).
- Lightweight logging for debugging runs.

4. Detailed Specifications

4.1 ReAct Loop

The agent must follow: **User Query** → **Thought** → **Action** → **Observation** → **Thought** → ⋯ → **Final Answer**.

Formatting contract (for parsing):

```
THOUGHT: <brief plan>
ACTION: <tool_name>{<valid JSON args>}
OBSERVATION: <tool result (trimmed)>
...
FINAL ANSWER: <concise answer>
```

4.2 Tool Specifications

Tool: `list_tables`

Description: Lists all tables in the database.

Parameters: None.

Returns: List of table names.

Tool: `describe_table`

Parameters: `table_name` (str).

Returns: Schema details (columns and types; row count if feasible).

Tool: `query_database`

Description: Executes a SELECT query (read-only).

Parameters: `query` (str). (You may also support `params` for prepared statements.)

Returns: Formatted result table or clear error message.

4.3 Safety, Validation, and Scope

- **Read-only:** Only `SELECT ... FROM ... [WHERE ...] [GROUP BY ...] [ORDER BY ...] [LIMIT ...]`.
- **Validation:** Reject non-SELECT statements. Whitelist identifiers (tables/columns) derived from `describe_table`.
- **Limits:** Enforce row caps (e.g., `LIMIT 100`) and reasonable timeouts.
- **Joins:** Optional. If included, keep to a single `INNER JOIN` or provide a pre-aggregated view in the DB.

4.4 Error Handling Requirements

Your agent must gracefully handle:

- **Invalid SQL:** Catch and surface syntax errors with a helpful suggestion.
- **Non-SELECT:** Refuse and explain the read-only constraint.
- **Parsing errors:** If LLM output deviates from the contract, request a properly formatted step (within the step budget).
- **Missing tables/columns:** Suggest using `list_tables` or `describe_table`.
- **Max step limit:** Return the best partial result and a summary of attempts.
- **LLM API errors:** Implement exponential backoff with a small retry budget.

5. Deliverables

- Code (single small package or single file):** Agent class, tools, prompts, and utilities; within 350–400 LOC (excl. comments/docstrings).
- README (1 page):** Purpose, how to run, tool list (arguments & returns), and a short example run with trace.
- Tests (3–5):** Cover schema discovery, simple aggregation, and one error-recovery case.
- Trace Logs:** A sample run with full THOUGHT/ACTION/OBSERVATION/FINAL_ANSWER.
- Reflection (≤ 300 words):** Design trade-offs, failure modes, and future directions.

6. Grading Rubric (100 points)

Criterion	Points
ReAct fidelity (clear Thought/Action/Observation, clean termination)	20
Tooling (design, validation, helpful error messages)	20
NL→SQL handling (intents, predicates, grouping/ordering, limits, safety)	25
Safety & Validation (read-only, identifier whitelisting, row/time caps)	15
Tests & Reproducibility (clear instructions, passes on provided DB)	10
Trace Quality (auditable, minimal yet sufficient)	10
Total	100

7. Submission Instructions

- Submit a zipped archive or repository link containing:
 - Source code in the form of a Github repository.
 - A detailed PDF file containing your report.
- Include a one-line command to run your agent on a sample query.
- **Deadline: November 10, 2025 (23:59 IST).**

8. Team & Integrity

- **Team size:** Up to **2** students.
- All work must be your own. Cite any external resources used (e.g., small parsing ideas or SQL patterns).
- You *may not* use agent frameworks; the agent, tools, and loop must be implemented **from scratch**.

Emphasis: This assignment is explicitly about building your own agent library *entirely from scratch*. You will design the ReAct loop, the tool registry, prompt scaffolding, SQL validation, and error handling yourself, within a concise code length.

Questions? Post on the course Discord.